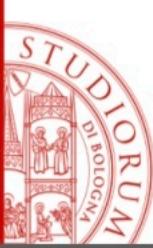


Controlli nascosti. Premi **ESC** per visualizzare i controlli.

Ignora

Databases

Relational Algebra and Relational Calculus



- Operations on the schema
 - DDL: data definition language
- Operations on data
 - DML: data manipulation language
 - Query instructions
 - Update instructions (to insert new data or modify current data)



Query languages for relational DBs

■ Declarative

They specify **what** are the results properties we want to obtain?

■ Procedural

They specify **how** the result is obtained?



Query Languages

- Relational Algebra: procedural
- Relational calculus: declarative
(theoretical, not implemented)
- SQL (Structured Query Language):
partially declarative (implemented)
- QBE (Query by Example):
declarative (implemented)



Relational Algebra

- Defined by a set of operators
 - over the relations
 - ...that produce relations as a result
 - ...and can be compositional



Relational Algebra Operators

- union, intersection, difference
- rename
- select
- project
- join (natural j., cartesian product, θ -join)



Set operators

- Relations are sets
- The results must also be sets
- we can use **union**, **intersection**, **difference** if the involved relations have the same schema



Union

Graduated

| Number | Name | Age |
|--------|-------|-----|
| 7274 | Rossi | 42 |
| 7432 | Neri | 54 |
| 9824 | Verdi | 45 |

Technicians

| Number | Name | Age |
|--------|-------|-----|
| 9297 | Neri | 33 |
| 7432 | Neri | 54 |
| 9824 | Verdi | 45 |

Graduated ∪ Technicians

| Number | Name | Age |
|--------|-------|-----|
| 7274 | Rossi | 42 |
| 7432 | Neri | 54 |
| 9824 | Verdi | 45 |
| 9297 | Neri | 33 |



Intersection

Graduated

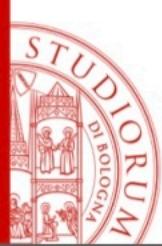
| Number | Name | Age |
|--------|-------|-----|
| 7274 | Rossi | 42 |
| 7432 | Neri | 54 |
| 9824 | Verdi | 45 |

Technicians

| Number | Name | Age |
|--------|-------|-----|
| 9297 | Neri | 33 |
| 7432 | Neri | 54 |
| 9824 | Verdi | 45 |

Graduated \cap Technicians

| Number | Name | Age |
|--------|-------|-----|
| 7432 | Neri | 54 |
| 9824 | Verdi | 45 |



Difference

Graduated

| Number | Name | Age |
|--------|-------|-----|
| 7274 | Rossi | 42 |
| 7432 | Neri | 54 |
| 9824 | Verdi | 45 |

Technicians

| Number | Name | Age |
|--------|-------|-----|
| 9297 | Neri | 33 |
| 7432 | Neri | 54 |
| 9824 | Verdi | 45 |

Graduated -
Technicians

| Number | Name | Age |
|--------|-------|-----|
| 7274 | Rossi | 42 |
| | | |
| | | |



A correct (but impossible) union

Fatherhood

| Father | Child |
|---------|--------|
| Adam | Abel |
| Adam | Cain |
| Abraham | Isacco |

Motherhood

| Mother | Child |
|--------|-------|
| Eve | Abel |
| Eve | Seth |
| Sarah | Isaac |

Fatherhood \cup Motherhood

??



Renaming

- Unary operator (only one argument)
- It produces a result that “changes the schema” while keeping the contained data unaltered.



Renaming: Syntax+Semantics

■ Syntax:

$\rho_{\text{NewName} \leftarrow \text{OldName}}$ (Relation)

■ Semantics:

Changes the attribute “OldName” into
“NewName”



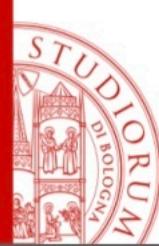
Renaming, examples

Fatherhood

| Father | Child |
|---------|-------|
| Adam | Abel |
| Adam | Cain |
| Abraham | Isaac |

$\rho_{\text{Parent} \leftarrow \text{Father}}$ (Fatherhood)

| Parent | Child |
|---------|-------|
| Adam | Abel |
| Adam | Cain |
| Abraham | Isaac |



Renaming, examples

Fatherhood

| Father | Child |
|---------|-------|
| Adam | Abel |
| Adam | Cain |
| Abraham | Isaac |

$\rho_{\text{Parent} \leftarrow \text{Father}}$ (Fatherhood)

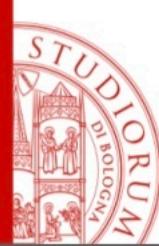
| Parent | Child |
|---------|-------|
| Adam | Abel |
| Adam | Cain |
| Abraham | Isaac |

Motherhood

| Mother | Child |
|--------|-------|
| Eve | Abele |
| Eve | Seth |
| Sarah | Isaac |

$\rho_{\text{Parent} \leftarrow \text{Mother}}$ (Motherhood)

| Parent | Child |
|--------|-------|
| Eve | Abel |
| Eve | Seth |
| Sarah | Isaac |



Renaming, examples

$\rho_{\text{Parent} \leftarrow \text{Father}}$ (Fatherhood)

| Parent | Child |
|--------|--------|
| Adamo | Abele |
| Adamo | Caino |
| Abramo | Isacco |

$\rho_{\text{Parent} \leftarrow \text{Mother}}$ (Motherhood)

| Parent | Child |
|--------|--------|
| Eva | Abele |
| Eva | Set |
| Sara | Isacco |

$\rho_{\text{Parent} \leftarrow \text{Father}}$ (Fatherhood) \cup
 $\rho_{\text{Parent} \leftarrow \text{Mother}}$ (Motherhood)

| Parent | Child |
|--------|--------|
| Adamo | Abele |
| Adamo | Caino |
| Abramo | Isacco |
| Eva | Abele |
| Eva | Set |
| Sara | Isacco |



Employee

| Surname | Office | Salary |
|---------|--------|--------|
| Rossi | Rome | 55 |
| Neri | Milan | 64 |

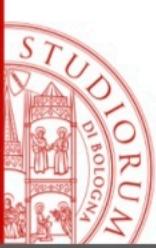
Worker

| Surname | Factory | Wage |
|---------|---------|------|
| Bruni | Paris | 45 |
| Verdi | Berlin | 55 |

$\rho_{\text{Office}, \text{Pay} \leftarrow \text{Office}, \text{Salary}}$ (Employee) \cup

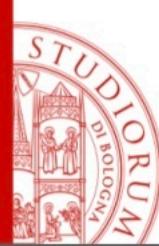
$\rho_{\text{Office}, \text{Pay} \leftarrow \text{Factory}, \text{Wage}}$ (Worker)

| Surname | Office | Pay |
|---------|--------|-----|
| Rossi | Rome | 55 |
| Neri | Milan | 64 |
| Bruni | Paris | 45 |
| Verdi | Berlin | 55 |



Selection

- Unary operator
- As a result:
 - The output has the same schema of the input
 - The output tuples are a subset of the input tuples
 - The output tuples satisfy a predicate



Selection, Syntax+Semantics

■ Syntax

$$\sigma_{\text{Predicate}} (\textit{Relation})$$

Predicate: Its interpretation is a boolean expression over the relations' tuples

■ Semantics

A subset of the relations that satisfy the given predicate



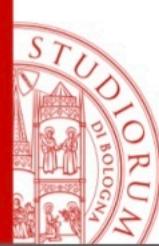
Selection, examples

Employee

| Number | Surname | Office | Salary |
|--------|---------|--------|--------|
| 7309 | Rossi | Rome | 55 |
| 5998 | Neri | Milan | 64 |
| 9553 | Milan | Milan | 44 |
| 5698 | Neri | Naples | 64 |

Return all the employees that:

- Earn more than 50
- Earn more than 50 and work in Milan
- Have the same surname as their city's office

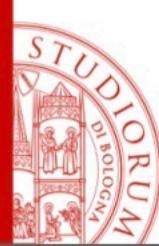


Selection, examples (1)

- ... Employees earning more than 50

$\sigma_{\text{Salary} > 50} (\text{Employee})$

| Number | Surname | Office | Salary |
|--------|---------|--------|--------|
| 7309 | Rossi | Rome | 55 |
| 5998 | Neri | Milan | 64 |
| 5698 | Neri | Naples | 64 |

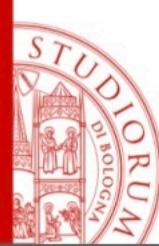


Selection, examples (2)

- ...Employees earning more than 50 and working in Milan

$\sigma_{\text{Salary} > 50 \text{ AND } \text{Office} = \text{'Milan'}} (\text{Employee})$

| Number | Surname | Office | Salary |
|--------|---------|--------|--------|
| 5998 | Neri | Milan | 64 |



Selection, examples (3)

- ...Employees having the same surname as their city's office

$$\sigma_{\text{Surname} = \text{Office}}(\text{Employee})$$

| Number | Surname | Office | Salary |
|--------|---------|--------|--------|
| 9553 | Milan | Milan | 44 |
| | | | |
| | | | |
| | | | |



Projection

- Unary operator
- As a result:
 - the output's schema is a subset of the inputs' schema
 - the output is formed using all the input's tuples



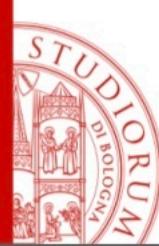
Projection, Syntax+Semantics

■ Syntax:

$$\pi_{\text{AttributeList}} (\text{Relation})$$

■ Semantics:

The results contains all the tuples in Relation, but only the attributes in **AttributeList**



Projection, examples

Employee

| Number | Surname | Office | Salary |
|--------|---------|--------|--------|
| 7309 | Neri | Naples | 55 |
| 5998 | Neri | Milan | 64 |
| 9553 | Rossi | Rome | 44 |
| 5698 | Rossi | Rome | 64 |

- For all the employees return
 - Number and Surname
 - Surname and Office



Projection, examples (1)

- Return the Employees' number and surname

π Number, Surname (Employee)

| Number | Surname |
|--------|---------|
| 7309 | Neri |
| 5998 | Neri |
| 9553 | Rossi |
| 5698 | Rossi |

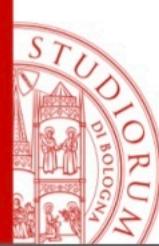


Projection, examples (2)

- Surname and Office for all the Employees

$\pi_{\text{Surname, Office}} (\text{Employee})$

| Surname | Office |
|---------|--------|
| Neri | Naples |
| Neri | Milan |
| Rossi | Rome |



Projections' cardinality

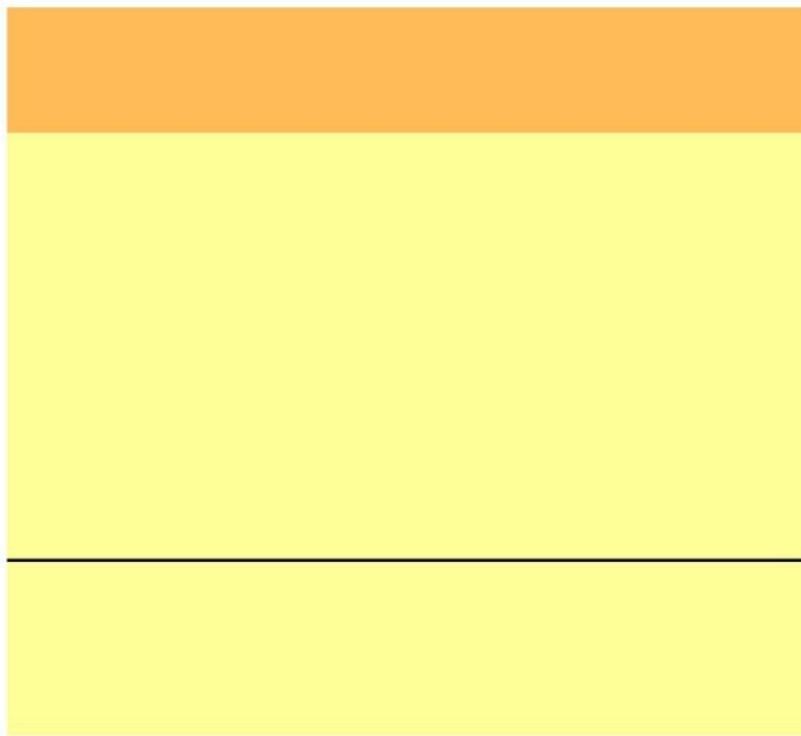
- A projection's output
 - contains at most the same number of tuples than the input's
 - could contain less tuples: restricting to a subset of the attributes, some tuples could be repeated! Repeated tuples are discarded in relations.
- if X is a superkey for R , then $\pi_X(R)$ contains the same number of tuples as R



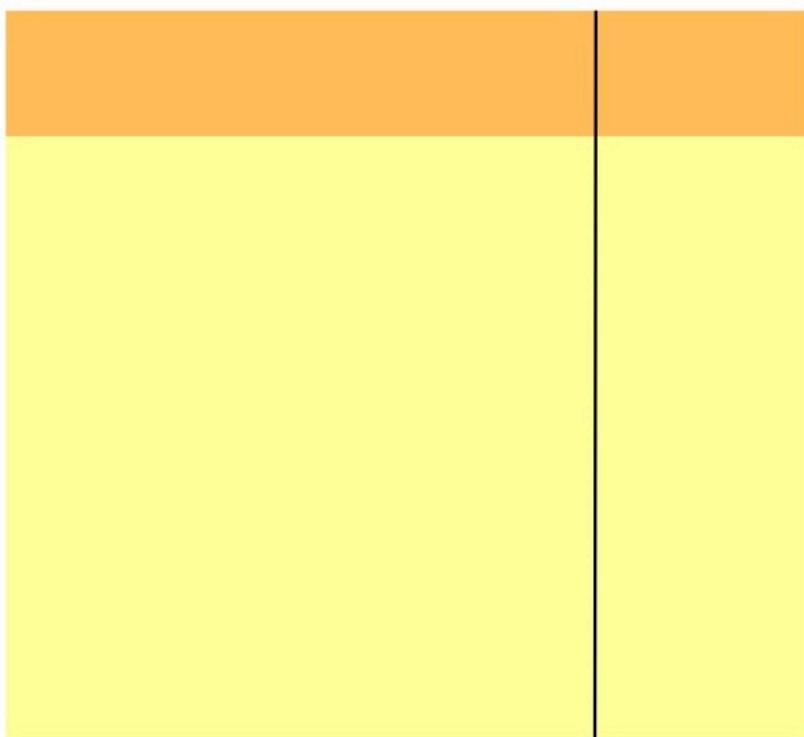
Selection and Projection (1)

- “orthogonal” operators
- **selections**: horizontal decomposition
- **projection**: vertical decomposition

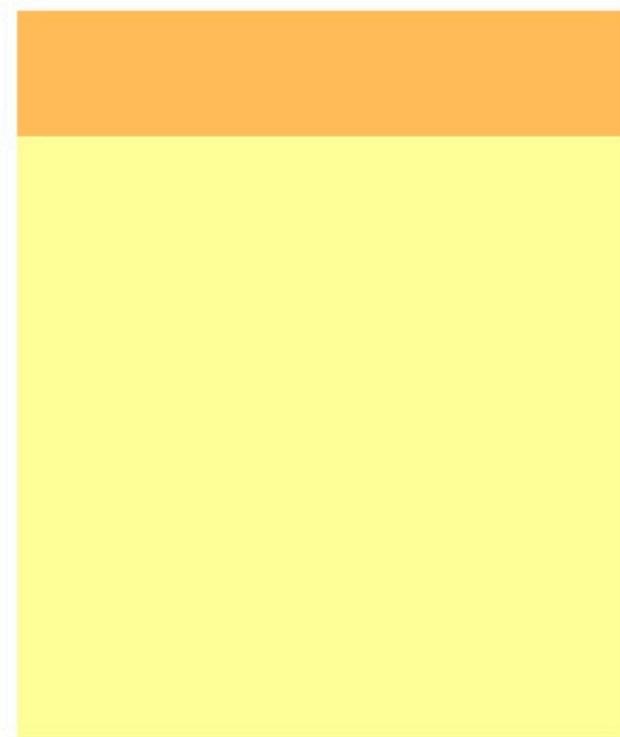
Selection and Projection (2)



selection



projection





Selection and Projection (3)

- By combining selection and projection, it is possible to extract only the interesting information from a relation.
- Let's see an example.

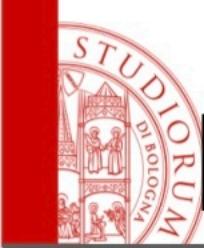


Selection and Projection (4)

- Return the number and the surname of the employees' having a salary greater than 50

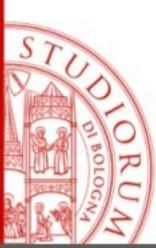
$$\pi_{\text{Number}, \text{Surname}} (\sigma_{\text{Salary} > 50} (\text{Employees}))$$

| Number | Surname |
|--------|---------|
| 7309 | Rossi |
| 5998 | Neri |
| 5698 | Neri |



Limits of the combo Selection+Projection

- By using such operators we could only extract informations from one single relation
- Hereby we cannot correlate neither tuples from different relations, nor different tuples from the same relation



Join

- The *join* is the most interesting operator of the relational algebra
- It allows the correlation between tuples in different relations



Exams during a public contest

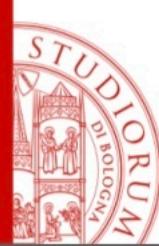
- Each exam is anonymous and a closed envelope, containing the candidate's surname, is associated to it
- Each exam and its related envelope have the same ID.



Data

| | | | |
|---|----|---|----------------|
| 1 | 25 | 1 | Jan Johansson |
| 2 | 13 | 2 | Jean B. Lully |
| 3 | 27 | 3 | Johann S. Bach |
| 4 | 28 | 4 | Giuseppe Verdi |

| | |
|----------------|----|
| Jan Johansson | 25 |
| Jean B. Lully | 13 |
| Johann S. Bach | 27 |
| Giuseppe Verdi | 28 |



Data with Schema

| ID | Grade |
|----|-------|
| 1 | 25 |
| 2 | 13 |
| 3 | 27 |
| 4 | 28 |

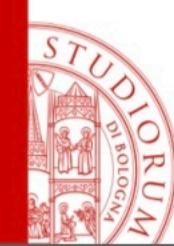
| ID | Candidate |
|----|----------------|
| 1 | Jan Johansson |
| 2 | Jean B. Lully |
| 3 | Johann S. Bach |
| 4 | Giuseppe Verdi |

| ID | Candidate | Grade |
|----|----------------|-------|
| 1 | Jan Johansson | 25 |
| 2 | Jean B. Lully | 13 |
| 3 | Johann S. Bach | 27 |
| 4 | Giuseppe Verdi | 28 |



Natural Join

- Binary operator (generalizable through associativity)
- Provides a result such that:
 - its schema is the union of the attributes of the two relations' schema
 - each tuple is produced combining two tuples: one from each of the two relations



Join, Syntax+Semantics

$R_1(X_1), R_2(X_2)$

$R_1 \bowtie R_2$ is a relation over $X_1 X_2$

$$R_1 \bowtie R_2 = \{ t \text{ on } X_1 X_2 \mid \\ \exists t_1 \in R_1 \text{ and } t_2 \in R_2 \text{ with } t[X_1] = t_1 \text{ and } t[X_2] = t_2 \}$$

Full Join

| Employee | Dept. | | Dept. | Chief |
|----------|-------|---|-------|-------|
| Rossi | A | ⊗ | A | Mori |
| Neri | B | | B | Bruni |
| Bianchi | B | | | |

| Employee | Dept. | Chief |
|----------|-------|-------|
| Rossi | A | Mori |
| Neri | B | Bruni |
| Bianchi | B | Bruni |

- **Full join:** each tuple contributes to the final result

A “not full” join

| Employee | Dept. | Dept. | Chief |
|----------|-------|-------|-------|
| Rossi | A | B | Mori |
| Neri | B | C | Bruni |
| Bianchi | B | | |

| Employee | Dept. | Chief |
|----------|-------|-------|
| Neri | B | Mori |
| Bianchi | B | Mori |



An “empty” join

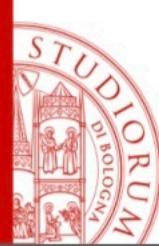
| Employee | Dept. | Dept. | Chief |
|----------|-------|-------|-------|
| Rossi | A | D | Mori |
| Neri | B | C | Bruni |
| Bianchi | B | | |

| Employee | Dept. | Chief |
|----------|-------|-------|
|----------|-------|-------|

A full join, having n×m tuples

| Employee | Dept. | | Dept. | Chief |
|----------|-------|---|-------|-------|
| Rossi | B | ⊗ | B | Mori |
| Neri | B | | B | Bruni |

| Employee | Dept. | Chief |
|----------|-------|-------|
| Rossi | B | Mori |
| Rossi | B | Bruni |
| Neri | B | Mori |
| Neri | B | Bruni |



Size of a join's result

- The result of the join between R_1 and R_2 has a number of tuples between zero and $|R_1| \times |R_2|$

$$0 \leq |R_1 \bowtie R_2| \leq |R_1| \times |R_2|$$

- If the join involves a key from R_2 , then the number of the resulting tuples is within 0 and $|R_1|$

$$0 \leq |R_1 \bowtie R_2| \leq |R_1|$$

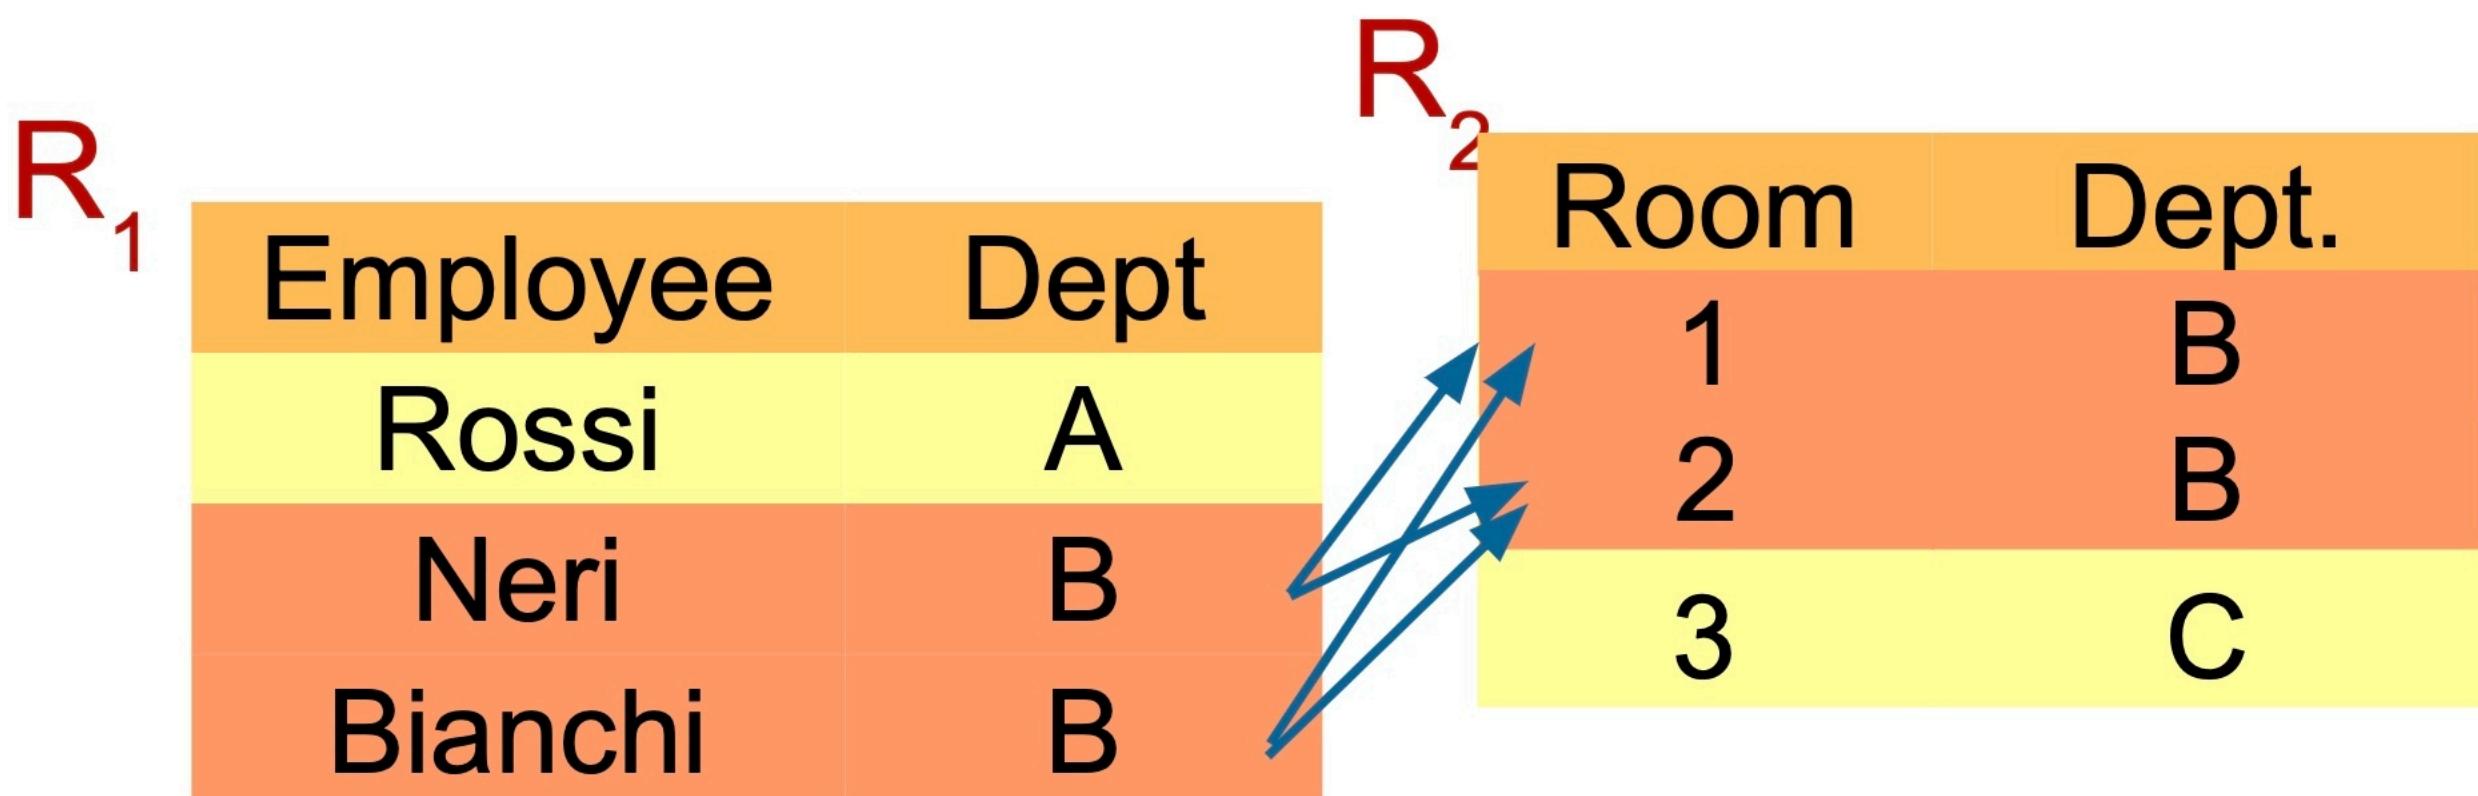
- If the join involves a key from R_2 and a Referential Integrity Constraint, the number of the tuples is $|R_1|$

$$|R_1 \bowtie R_2| = |R_1|$$

Size of a join's result, example (1)

- $|R_1|$ and $|R_2|$ have a size of 3
- The size of the join between R_1 and R_2 is at most 9 (3×3)

| R_1 | | | R_2 | | |
|-------|----------|------|-------|------|-------|
| | Employee | Dept | | Room | Dept. |
| | Rossi | A | | 1 | B |
| | Neri | B | | 2 | B |
| | Bianchi | B | | 3 | C |



$$|R_1 \bowtie R_2| = 4$$

Size of a join's result, example (2)

- The size of the join between R_1 and R_2 involves R_2 's key.
- Hence its size is between 0 and $|R_1|$
- Since the key's value are unique, for each tuple of R_2 can match more tuples of R_1 and not vice versa

| R_1 | Employee | Dept. | R_2 | Dept. | Chief |
|-------|----------|-------|-------|-------|-------|
| | Rossi | A | | B | Mori |
| | Neri | B | | C | Bruni |
| | Bianchi | B | | | |

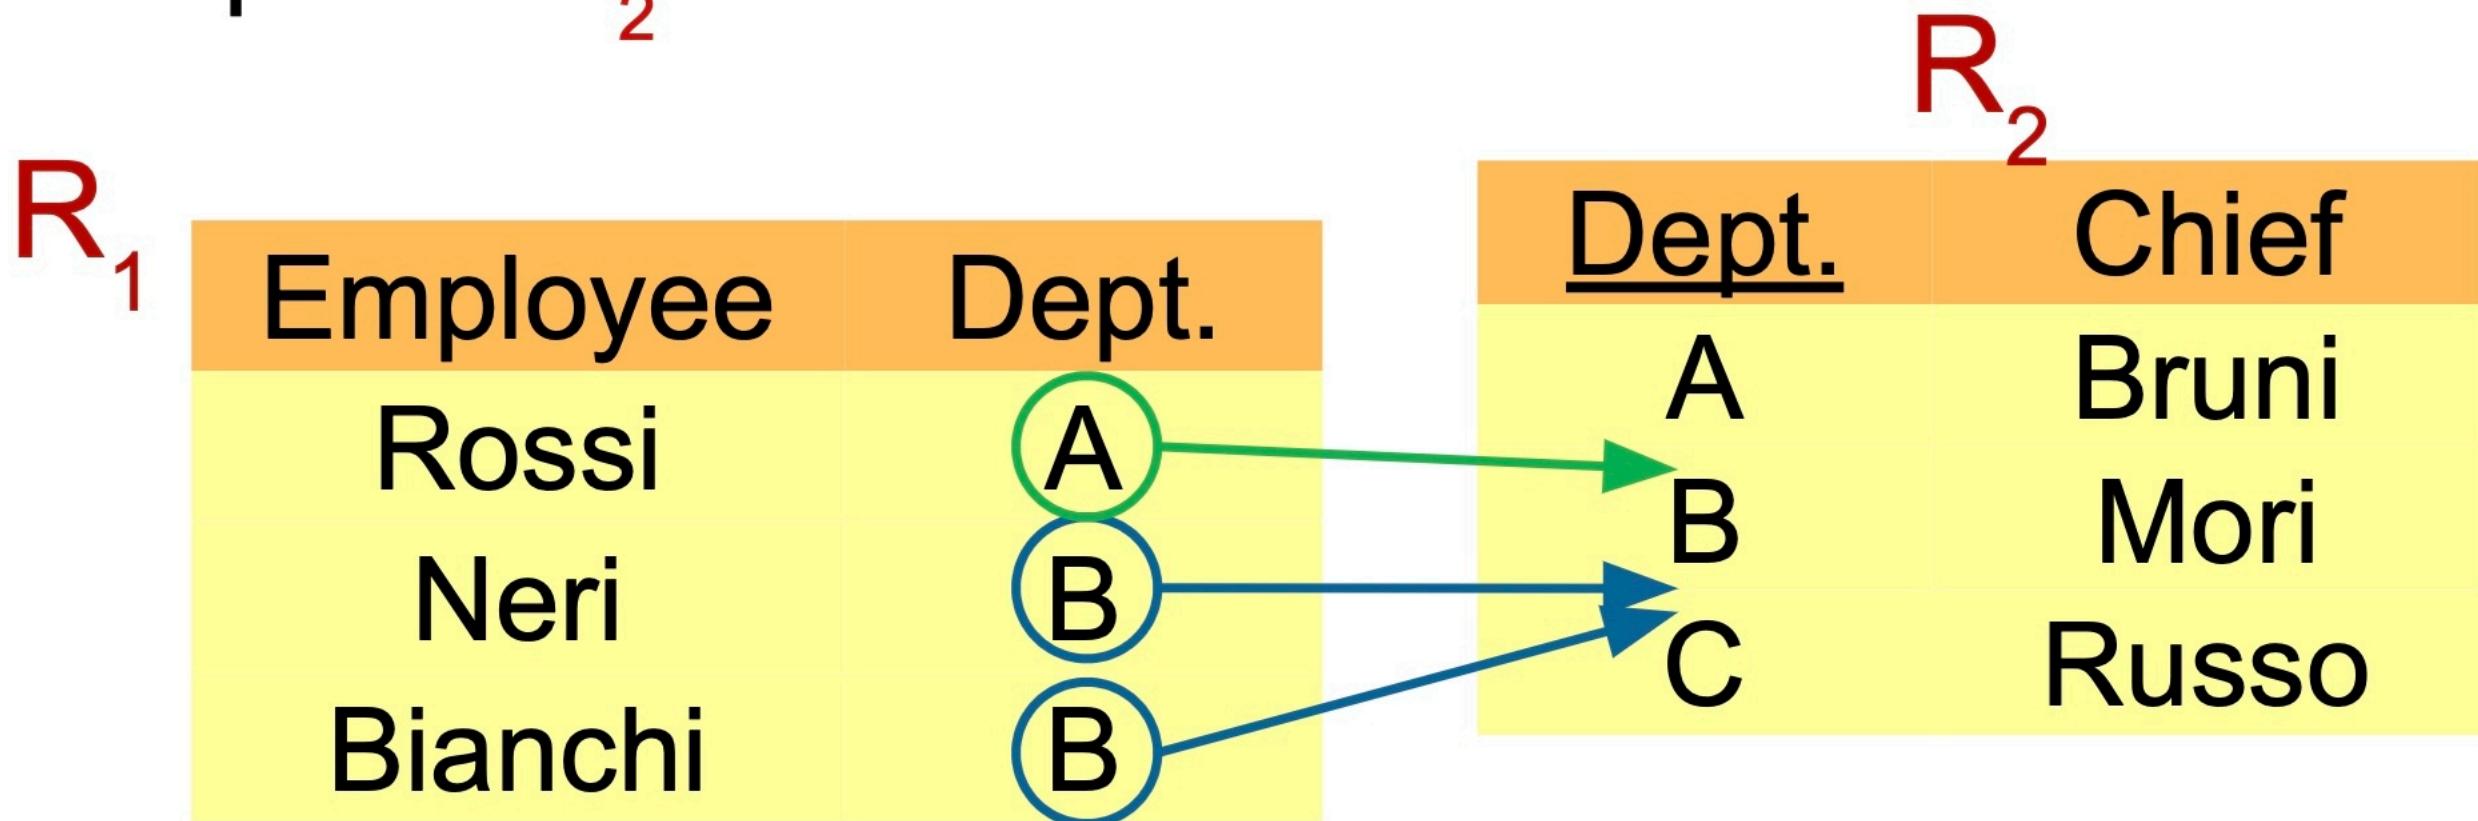
The diagram illustrates a join operation between two relations, R_1 and R_2 . Relation R_1 has columns 'Employee' and 'Dept.', with rows ('Rossi', 'A'), ('Neri', 'B'), and ('Bianchi', 'B'). Relation R_2 has columns 'Dept.' and 'Chief', with rows ('B', 'Mori') and ('C', 'Bruni'). A blue arrow points from the 'B' in the 'Dept.' column of R_1 to the 'B' in the 'Dept.' column of R_2 , indicating a match. Another blue arrow points from the second 'B' in the 'Dept.' column of R_1 to the 'C' in the 'Dept.' column of R_2 , indicating another match. The third row in R_1 ('Bianchi', 'B') has no corresponding row in R_2 , demonstrating that not every tuple in R_1 matches a tuple in R_2 .

Size of a join's result, example (3)

- If there exists a referential integrity constraint between Dept. (in R_1) and Dept. (key in R_2):

$$|R_1 \bowtie R_2| = |R_1|$$

each tuple in R_1 is associated to at least one tuple in R_2

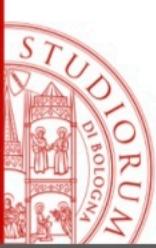


Join, critical issues

| Employee | Dept. | Dept. | Chief |
|----------|-------|-------|-------|
| Rossi | A | B | Mori |
| Neri | B | C | Bruni |
| Bianchi | B | | |

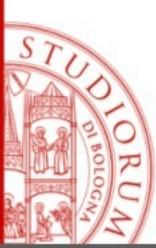
| Employee | Dept. | Chief. |
|----------|-------|--------|
| Neri | B | Mori |
| Bianchi | B | Mori |

Some tuples do not contribute to the final result: they are “left out”.



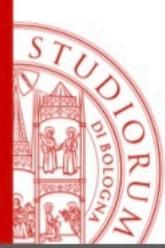
Outer Join

- The **outer** join fills with **NULL** values the tuples that are normally discarded from an “**inner**” join
- There are three types of “outer joins”: left, right, and full.



Outer Join

- **left (⊗)**: keeps all the tuples from the first operand, even if there are no right operand matching tuples. The latter are replaced by NULL values.
- **right (⊗)**: $A \otimes B$ is defined as $B \otimes A$
- **full (⊛)**: Is defined as the union of the former operands



Left Outer Join

Employee

| Employee | Dept. |
|----------|-------|
| Rossi | A |
| Neri | B |
| Bianchi | B |

Department

| Dept. | Chief. |
|-------|--------|
| B | Mori |
| C | Bruni |

Employee \bowtie Department

| Employee | Dept. | Chief |
|----------|-------|-------|
| Neri | B | Mori |
| Bianchi | B | Mori |
| Rossi | A | NULL |



Right Outer Join

Employee

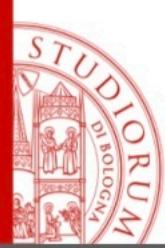
| Employee | Dept. |
|----------|-------|
| Rossi | A |
| Neri | B |
| Bianchi | B |

Department

| Dept. | Chief |
|-------|-------|
| B | Mori |
| C | Bruni |

Employee \bowtie Department

| Employee | Dept. | Chief. |
|----------|-------|--------|
| Neri | B | Mori |
| Bianchi | B | Mori |
| NULL | C | Bruni |



Full Outer Join

Employee

| Employee | Dept. |
|----------|-------|
| Rossi | A |
| Neri | B |
| Bianchi | B |

Department

| Dept. | Chief |
|-------|-------|
| B | Mori |
| C | Bruni |

Employee \bowtie Department

| Employee | Dept. | Chief |
|----------|-------|-------|
| Neri | B | Mori |
| Bianchi | B | Mori |
| Rossi | A | NULL |
| NULL | C | Bruni |



Employee

| Employee | Dept. |
|----------|-------|
| Rossi | A |
| Neri | B |
| Bianchi | B |

Department

| Code | Chief |
|------|-------|
| A | Mori |
| B | Bruni |

Employee \bowtie Department

| Employee | Dept. | Code | Chief |
|----------|-------|------|-------|
| Rossi | A | A | Mori |
| Rossi | A | B | Bruni |
| Neri | B | A | Mori |
| Neri | B | B | Bruni |
| Bianchi | B | A | Mori |
| Bianchi | B | B | Bruni |



Cartesian Product

- It's a Natural Join over two relations with no common attributes.
- The result's size is the product of the operands' size.
- In practice, cartesian product is useful only if followed by a selection:

$$\sigma_{\text{condition}} (R1 \bowtie R2)$$



Theta-join: why?

- The natural join assumes a specific meaning when it is followed by a selection operation:

$$\sigma_{\text{condition}}(R_1 \bowtie R_2)$$

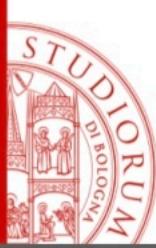
- Such operation is defined as theta-join

$$R_1 \bowtie_{\text{condition}} R_2$$



Theta-join (2)

- The predicate **condition** is usually defined as a conjunction (**AND**) of atoms expressing binary relations $A_1 \mathcal{R} A_2$, where \mathcal{R} is a comparison operator ($=, >, <, \dots$)



Equi-join

- A theta-join is called “equi-join” iff all the \mathcal{R} atoms in theta are equivalence relations.

Note that the equi-join is used much more often than the most general theta-join: it allows to specify on which attribute to join without using rename operations.



Employee

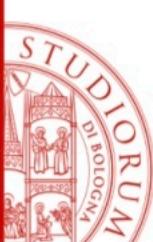
| Employee | Dept. |
|----------|-------|
| Rossi | A |
| Neri | B |
| Bianchi | B |

Department

| Code | Chief |
|------|-------|
| A | Mori |
| B | Bruni |

Employee $\bowtie_{\text{Dept.}=\text{Code}}$ Department

| Employee | Dept. | Code | Chief |
|----------|-------|------|-------|
| Rossi | A | A | Mori |
| Neri | B | B | Bruni |
| Bianchi | B | B | Bruni |



Please Note

- The **equi-join** provides a similar result as the join between **Employee** and **Department** where “Code” is renamed as “Dept.”

$\text{Employee} \bowtie p_{\text{Dept.} \leftarrow \text{Code}} (\text{Department})$

| Employee | Dept. | Chief |
|----------|-------|-------|
| Rossi | A | Mori |
| Neri | B | Bruni |
| Bianchi | B | Bruni |



Natural Join and EquiJoin

- After renaming we obtain the two following schemas, which can be joined with natural join. We then proceed to selection and projection.

Employee

| Employee | Dept. |
|----------|-------|
| | |

Department

| Dept. | Chief |
|-------|-------|
| | |

Employee \bowtie Department

$$\begin{array}{l} \pi_{\text{Employee}, \text{Dept.}, \text{Chief}} \\ (\sigma_{\text{Department} = \text{Code}} \\ (\rho_{\text{Code} \leftarrow \text{Dept.}} (\text{Department}))) \end{array}$$



Examples

Employee

| Code | Name | Age | Wage |
|------|---------|-----|------|
| 7309 | Rossi | 34 | 45 |
| 5998 | Bianchi | 37 | 38 |
| 9553 | Neri | 42 | 35 |
| 5698 | Bruni | 43 | 42 |
| 4076 | Mori | 45 | 50 |
| 8123 | LUPI | 46 | 60 |

Supervisor

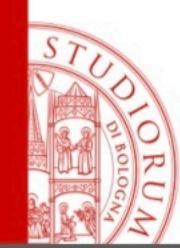
| Employee | Chief |
|----------|-------|
| 7309 | 5698 |
| 5998 | 5698 |
| 9553 | 4076 |
| 5698 | 4076 |
| 4076 | 8123 |



Relational Algebra: Examples (1)

- Return the employees' number, name, age and wage earning more than 40

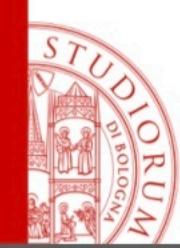
$$\sigma_{\text{Wage} > 40}(\text{Employee})$$



Relational Algebra: Examples (2)

- Return the employees' number, name and age earning more than 40

$$\pi_{\text{Number}, \text{Name}, \text{Age}} (\sigma_{\text{Wage} > 40}(\text{Employee}))$$



Relational Algebra: Examples (3)

- Return the employees' chiefs earning more than 40

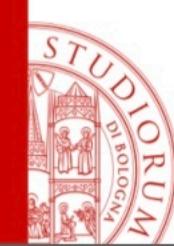
$$\pi_{\text{Chief}}(\text{Supervisor} \bowtie_{\text{Employee}=\text{Number}} (\sigma_{\text{Salary}>40}(\text{Employee})))$$



Relational Algebra: Examples (4)

- Return the chief's name and wage having employees earning more than 40

$$\begin{aligned} & \pi_{\text{Name}, \text{Wage}} (\text{Employee} \bowtie \\ & \quad \text{Number} = \text{Chief} \\ & \quad \pi_{\text{Chief}} (\text{Supervisor} \bowtie \\ & \quad \quad \text{Employee} = \text{Number} \\ & \quad \quad \sigma_{\text{Wage} > 40} (\text{Employee}) \\ &) \\ &) \end{aligned}$$



Relational Algebra: Examples (5)

- Return the employees having a wage greater than their chief's, along with number, name and wage of both the employee and his chief

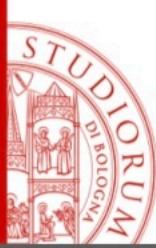
$$\begin{aligned} & \Pi_{\text{Code}, \text{Name}, \text{Wage}, \text{NoC}, \text{NameC}, \text{WageC}} (\\ & \quad \sigma_{\text{Salary} > \text{WageC}} (\\ & \quad \rho_{\text{NoC}, \text{NameC}, \text{WageC}, \text{AgeC} \leftarrow \text{Number}, \text{Name}, \text{Wage}, \text{Age}} (\text{Employee}) \\ & \quad \bowtie_{\text{NoC} = \text{Chief}} \\ & \quad (\text{Supervisor} \bowtie_{\text{Employee} = \text{Number}} \text{Employee}) \\ &) \\ &) \end{aligned}$$



Relational Algebra: Examples (6)

- Return the chiefs' number having all their employees earning more than 40

$$\begin{array}{l} \pi_{\text{Chief}}(\text{Supervision}) - \\ \pi_{\text{Chief}}(\text{Supervision} \\ \quad \bowtie \text{Employee} = \text{Number} \\ \quad \sigma_{\text{Wage} \leq 40}(\text{Employee}) \\ \end{array}$$



Equivalent Relational Algebra Expressions

- Two Relational Algebra expressions E_1 and E_2 are **equivalent** if they provide the same result independently from the database's state.
- Equivalence rules are very important because modern DBMSs try to rewrite the queries into equivalent (but more efficient, i.e. less time consuming) expressions



Equivalent Relational Algebra Expressions (2)

1. Combining a cascade of selections:

$$\sigma_{A>20} (\sigma_{B='Jones'} (R)) = \sigma_{A>20 \text{ AND } B='Jones'} (R)$$

2. Projection's idempotency (X and Y belong to R's schema)

$$\pi_X(R) = \pi_X(\pi_{XY}(R))$$



Equivalent Relational Algebra Expressions (3)

3. Push selections (A belongs to R_2 's schema)

$$\sigma_{A=10}(R_1 \bowtie R_2) = R_1 \bowtie (\sigma_{A=10}(R_2))$$

It drastically reduces the intermediate result's size (therefore also the operation cost).



Equivalent Relational Algebra Expressions (4)

4. Pushing projections

$$\pi_{X_1 Y_2}(R_1 \bowtie R_2) = R_1 \bowtie \pi_{Y_2}(R_2)$$

- R_1 has schema X_1 , R_2 has schema X_2
- If $Y_2 \subseteq (X_1 \cap X_2)$, the attributes in $X_2 - Y_2$ are not involved in the join, and hence the equivalence is valid.



Equivalent Relational Algebra Expressions (5)

5. Using the theta-join definition

a. $\sigma_{A=10} (R_1 \bowtie R_2) = R_1 \bowtie_{A=10} R_2$

b. $\sigma_{A=10 \text{ AND } B='Jones'} (R_1 \bowtie R_2) =$
 $\sigma_{A=10} (R_1 \bowtie_{B='Jones'} R_2) =$
 $\sigma_{B='Jones'} (R_1 \bowtie_{A=10} R_2)$



Equivalent Relational Algebra Expressions (6)

■ Some distributive properties

$$6. \sigma_{A=10}(R_1 \cup R_2) = \sigma_{A=10}(R_1) \cup \sigma_{A=10}(R_2)$$

$$7. \sigma_{A=10}(R_1 - R_2) = \sigma_{A=10}(R_1) - \sigma_{A=10}(R_2)$$

$$8. \pi_x(R_1 \cup R_2) = \pi_x(R_1) \cup \pi_x(R_2)$$

- Please note that projection is not distributive over difference



Equivalent Relational Algebra Expressions (7)

■ Alcune proprietà degli insiemi e della selezione

$$9. \sigma_{A=10 \text{ OR } B='Jones'}(R) = \sigma_{A=10}(R) \cup \sigma_{B='Jones'}(R)$$

$$10. \sigma_{A=10 \text{ AND } B='Jones'}(R) = \sigma_{A=10}(R) \cap \sigma_{B='Jones'}(R)$$

$$11. \sigma_{A=10 \text{ AND } \neg B='Jones'}(R) = \sigma_{A=10}(R) - \sigma_{B='Jones'}(R)$$

12. Commutativity for the union operator

$$R \bowtie (R_1 \cup R_2) = (R \bowtie R_1) \cup (R \bowtie R_2)$$



Example

Employee

| Number | Name | Age | Wage |
|--------|---------|-----|------|
| 7309 | Rossi | 34 | 45 |
| 5998 | Bianchi | 37 | 38 |
| 9553 | Neri | 42 | 35 |
| 5698 | Bruni | 43 | 42 |
| 4076 | Mori | 45 | 50 |
| 8123 | LUPI | 46 | 60 |

Supervisor

| Employee | Chief |
|----------|-------|
| 7309 | 5698 |
| 5998 | 5698 |
| 9553 | 4076 |
| 5698 | 4076 |
| 4076 | 8123 |



Example (1)

- Provide all the Chief's numbers being less than 30 years old.

$$\Pi_{\text{Chief}} (\sigma_{\text{No}=\text{Emp} \text{ AND } \text{Age}<30} (\text{Employee} \bowtie \text{Supervisor}))$$



Example (2)

- Rule №1, split the selection

$\Pi_{\text{Chief}}(\sigma_{\text{No}=\text{Emp}}(\sigma_{\text{Age}<30}(\text{Employee} \bowtie \text{Supervisor})))$

- Rule №5 for the $\text{No}=\text{Emp}$ selection and №3 for the $\text{age}<30$ selection

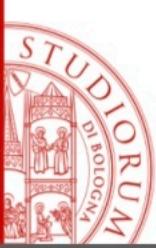
$\Pi_{\text{Chief}}(\sigma_{\text{Age}<30}(\text{Employee}) \bowtie_{\text{No}=\text{Emp}} \text{Supervisor})$



Example (3)

- Rule №4, delete all the unused attributes

$$\begin{aligned} & \pi_{\text{Chief}}(\pi_{\text{No}}(\sigma_{\text{Age} < 30}(\text{Employee})) \\ & \quad \bowtie_{\text{No} = \text{Emp}} \\ & \quad \text{Supervisor} \\ &) \\ &) \end{aligned}$$



Please Note

- You don't have to learn to write efficient relational algebra queries in this course (it's better having correct and clear answers)
- That's because modern DBMSs already do that work for you.



Selection with NULL values

$\sigma_{Age > 40} (\text{Employee})$

Employee

| Number | Surname | Agency | Age |
|--------|---------|--------|------|
| 7309 | Rossi | Rome | 32 |
| 5998 | Neri | Milan | 45 |
| 9553 | Bruni | Milan | NULL |

- No NULL value satisfies this specific atomic condition



Non Desiderata

- Selections are evaluated separately:
 $\sigma_{Age > 30}(\text{People}) \cup \sigma_{Age \leq 30}(\text{People}) \neq \text{People}$
- Atomic conditions are evaluated separately
 $\sigma_{Age > 30 \text{ OR } Age \leq 30}(\text{People}) \neq \text{People}$



NULL Valued selections (1)

- The following atomic condition is satisfied only by non-NULL values: $\sigma_{\text{Age} > 40}(\text{People})$
- Specific atomic statements are used for referring to NULL values:

IS NULL

IS NOT NULL
- We could even define a three-valued logic (based upon true, false, unknown), but it is not necessary.



NULL Valued selections (2)

Hereby:

$$\begin{aligned} \sigma_{Age > 30}(\text{People}) \cup \sigma_{Age \leq 30}(\text{People}) \\ \cup \sigma_{Age \text{ IS NULL}}(\text{People}) \\ = \\ \sigma_{Age > 30 \text{ OR } Age \leq 30 \text{ OR } Age \text{ IS NULL}}(\text{People}) \\ = \\ \text{People} \end{aligned}$$



NULL Valued selections (3)

σ (Age > 40) OR (Age IS NULL) (Employee)

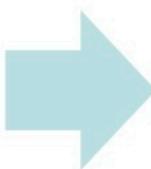
Employee

| Number | Surname | Agency | Age |
|--------|---------|--------|------|
| 5998 | Neri | Milan | 45 |
| 9553 | Bruni | Milan | NULL |

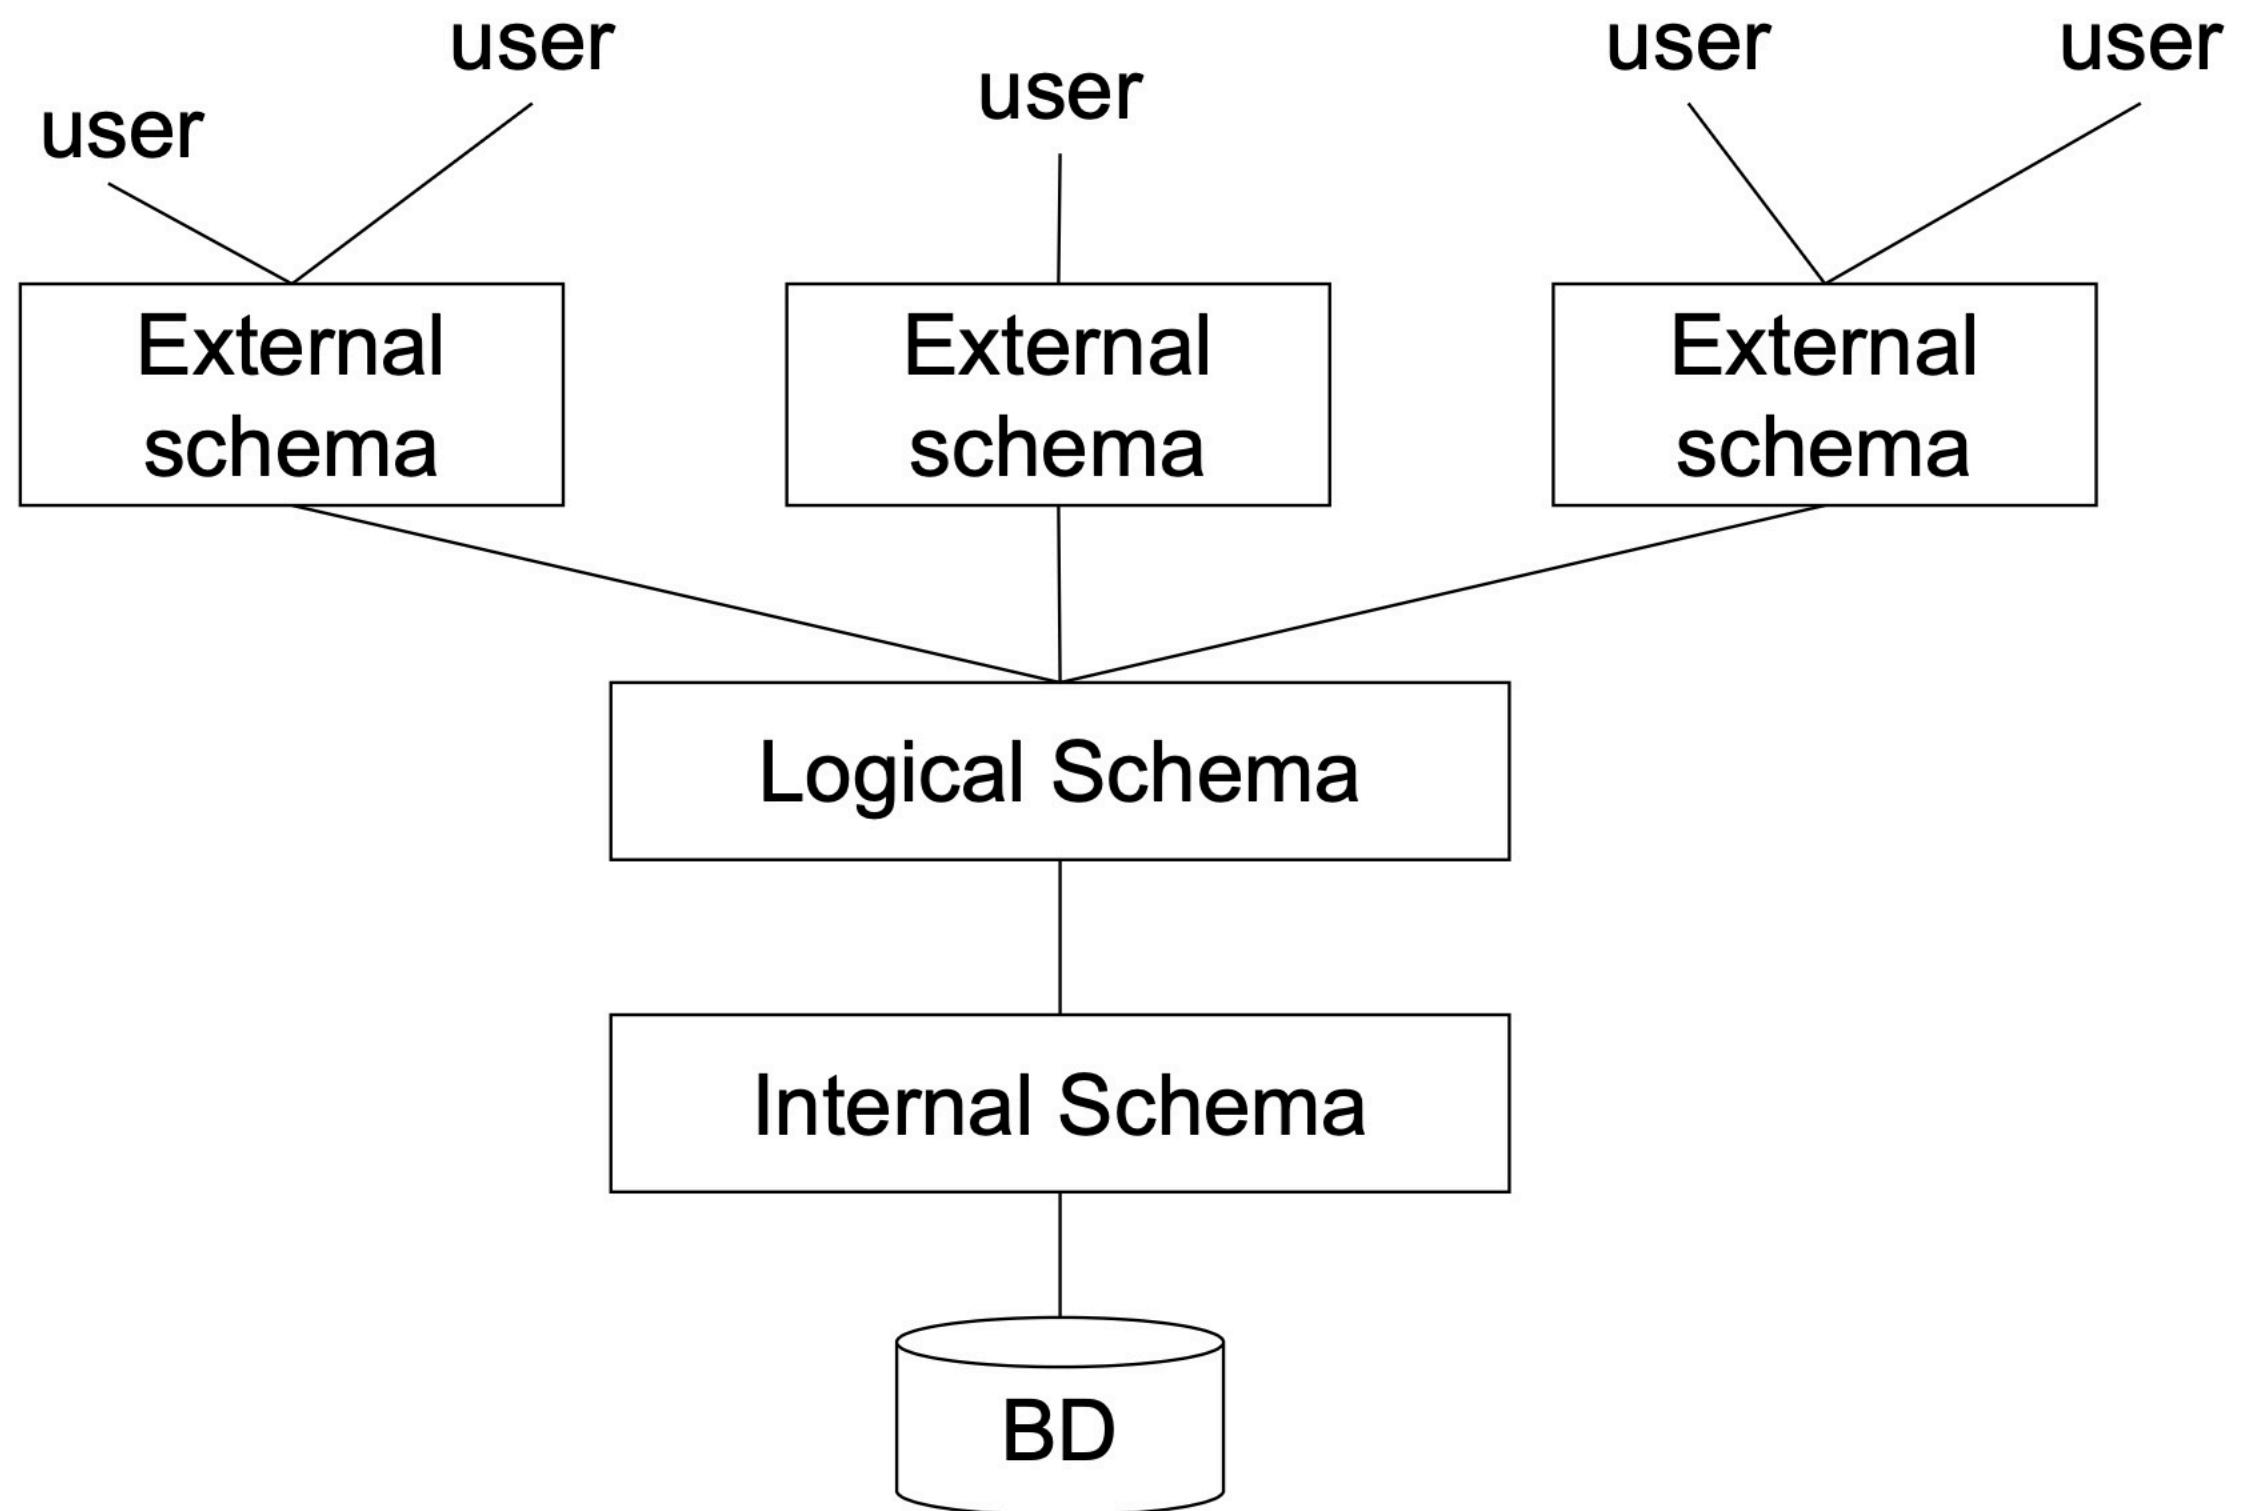


Views

- Different representations for the same data
- **Derived tables:** such relations are created from queries
- **Base tables:** original content
- Derived relations could be formed from other derived relations, but...



Standard three-tiered ANSI/SPARC architecture



Views, Example

Concern

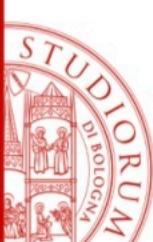
| Employee | Dept. |
|----------|-------|
| Rossi | A |
| Neri | B |
| Bianchi | B |

Manage

| Dept. | Chief |
|-------|-------|
| A | Mori |
| B | Bruni |

- Create a view:

Supervisor = $\pi_{\text{Employee}, \text{Chief}} (\text{Concern} \bowtie \text{Manage})$



Using Views

Views could be used in two fashions:

- Materialized views
- Virtual relations (or views)



View Materialization

- A temporary or permanent view table is stored physically
 - Pros:
 - Data is promptly available for further queries
 - Cons:
 - Data redundancy
 - Updates are slowed down
 - DBMS rarely support them



Virtual relations (1)

- Virtual relations (or views):
 - All the DBMS support them
 - The view query is transformed into a query on the underlying database



Virtual relations (2)

- The view's name is replaced by its associated query:

$\sigma_{\text{Chief}='Leoni'}$ (Supervisor)

Is run as:

$\sigma_{\text{Chief}='Leoni'}(\pi_{\text{Employee, Chief}}$ (Concern \bowtie Manage))

Views, reasons

- *External schema*: each user cannot only see:
 - both what (s)he is interested on and in the way (s)he likes it, with no further distractions
 - what (s)he is allowed to see
- *Tool in programming*:
 - we could simplify the writing of complex queries when sub-expressions are repeated
- Using already-existent software over pre-defined schemas

But:

- Views do not affect queries' efficiency



Views as a programming tool

- Return the employees having Jones's Chief

- Without views:

$$\begin{aligned} \pi_{\text{Employee}} & ((\text{Concern} \bowtie \text{Manage}) \bowtie \\ & \rho_{\text{ImpR,RepR} \leftarrow \text{Emp,Dept}} (\sigma_{\text{Employee}=\text{'Jones'}} (\text{Concern} \bowtie \\ & \text{Manage})) \\) \end{aligned}$$

- Using views:

$$\begin{aligned} \pi_{\text{Employee}} & (\text{Supervisor} \bowtie \\ & \rho_{\text{ImpR} \leftarrow \text{Imp}} (\sigma_{\text{Employee}=\text{'Jones'}} (\text{Supervisor})) \\) \end{aligned}$$


Updating views

Concern

| Employee | Dept |
|----------|------|
| Rossi | A |
| Neri | B |
| Verdi | A |

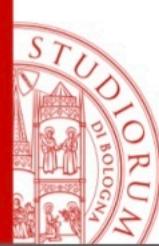
Manage

| Dept | Chief |
|------|-------|
| A | Mori |
| B | Bruni |
| C | Bruni |

Supervisor

| Employee | Chief |
|----------|-------|
| Rossi | Mori |
| Neri | Bruni |
| Verdi | Mori |

How could we update the data such that Bruni is Lupi's chief and that Falchi is Belli's chief?



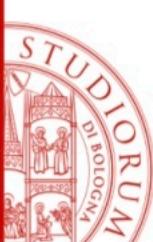
Incremental Updates

- To “update a view” means to change the base tables such that the updated view reflects the update
- To each update over the view must correspond to the one over the tables.
- Such update could be not unambiguous!
- Only a few possible updates are allowed on views



An alternative notation for joins (1)

- Please note: such approach is usually adopted for SQL implementations
- We ignore the *Natural Join* (we do not implicitly assume conditions over attributes with the same name)
- We disambiguate the attributes with the same name over different relations using the *Relation.Attribute* syntax
- We give relations a new name by creating views (and we rename attributes only when it's needed for the union operation).



Conventions and notations (1)

- Find the employees earning more than their chiefs, showing the number, name and wage of both employee and chief.

$$\begin{array}{l} \Pi_{\text{No}, \text{Name}, \text{Wage}, \text{NoC}, \text{NameC}, \text{WageC}} (\\ \sigma_{\text{Wage} > \text{WageC}} (\\ \rho_{\text{NoC}, \text{NameC}, \text{WageC}, \text{AgeC} \leftarrow \text{No}, \text{Name}, \text{Wage}, \text{Age}} (\text{Emp}) \\ \bowtie \\ \text{NoC} = \text{Chief} \\ (\text{Sup} \bowtie_{\text{Emp} = \text{No}} \text{Emp}) \\) \\) \end{array}$$



Conventions and notations (2)

- Assign Emp to Chief (View)
 $\text{Chief} := \text{Emp}$
- We use relations' name as prefix to differentiate between shared attributes
e.g. Emp.Wage and Chief.Wage

$$\begin{array}{l} \pi_{\text{Emp.No, Emp.Name, Emp.Wage, Chief.No, Chief.Name, Chief.Wage}} \\ \sigma_{\text{Emp.Wage} > \text{Chief.Wage}} \\ \text{Chief} \bowtie_{\text{Chief.No} = \text{Chief}} (\text{Sup} \bowtie_{\text{Emp} = \text{Emp.No}} \text{Emp}) \\) \\) \end{array}$$



Relational Calculi

- A family of declarative languages based on First Order Logic
- Two different definitions:
 - Domain Relational Calculus
 - Tuple Relational Calculus with Range Declarations



Domain Relational Calculus

- An expression is in the form:

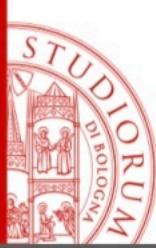
$$\{ A_1: x_1, \dots, A_k: x_k \mid f \}$$

- f is a formula (with boolean operators and quantifiers)
- $A_1: x_1, \dots, A_k: x_k$ target list:
 - A_1, \dots, A_k different attributes (can be in different databases)
 - x_1, \dots, x_k different variables
- Semantics: the result is a relation over A_1, \dots, A_k containing tuples of values for x_1, \dots, x_k satisfying the formula f



Comments

- Differences with Predicate Logic (for those who know it):
 - “predicate” symbols:
 - represent relations within the database
 - "standard" built-in predicates ($=$, $>$, ...)
 - There are no function symbols
 - The most interesting are “unbounded” predicates
 - Non-positional notation



Example

Employee(Number, Name, Age, Salary)

Supervisor(Chief, Employee)



Example (0a)

- Return the number, name, age and salary of the employees earning more than 40

$$\sigma_{\text{Salary} > 40}(\text{Employee})$$

{ Number: m , Name: n , Age: a , Salary: s |
Employee(Number: m , Name: n , Age: a , Salary: s)
 $\wedge s > 40$ }



Example (0b)

- Return the employees' number, name and age for all

$$\pi_{\text{Number}, \text{Name}, \text{Age}}(\text{Employee})$$

{ Number: m , Name: n , Age: a |
 $\exists s \text{ (Employee(Number: } m, \text{ Name: } n, \text{ Age: } a, \text{ Salary: } s\text{)}}$

{ Number: m , Name: n , Age: a |
Employee(Number: m , Name: n , Age: a , Salary: s)}



Example (1)

- Return the employees' number, name and age earning more than 40

$$\pi_{\text{Number, Name, Age}}(\sigma_{\text{Salary} > 40}(\text{Employee}))$$

{ Number: m , Name: n , Age: a |
Employee(Number: m , Name: n , Age: a , Salary:
 $s) \wedge s > 40 \}$



Example (2)

- Return the number identifying the employees' chiefs earning more than 40

$$\pi_{\text{Chief}}(\text{Supervisor} \bowtie_{\text{Employee}=\text{Number}} \sigma_{\text{Salary}>40}(\text{Employee}))$$

{ Chief: c | Supervisor(Chief: c , Employee: e) \wedge
Employee(Number: e , Name: n , Age: a , Salary: s) $\wedge s >$
40 }



Example (3)

- Return the chiefs' name and salary having employees earning more than 40

$$\pi_{\text{NameC}, \text{WageC}}(\rho_{\text{NoC}, \text{NameC}, \text{WageC}, \text{AgeC} \leftarrow \text{No}, \text{Name}, \text{Wage}, \text{Age}}(\text{Employee}) \\ \bowtie_{\text{NoC} = \text{Chief}} \\ (\text{Supervisor} \bowtie_{\text{Employee} = \text{Number}} \sigma_{\text{Salary} > 40}(\text{Employee})))$$

{ NameC: nc, WageC: sc |
Employee(Number: m, Name: n, Age: a, Salary: s) \wedge s >
40 \wedge Supervisor(Chief:c, Employee:m) \wedge
Employee(Number:c, Name:nc, Age:ac, Salary: sc) }



Example (4)

- Return the employees earning more money than their boss; for both such employees and chiefs return the number, name and salary

$$\begin{array}{c} \pi_{\text{No}, \text{Name}, \text{Wage}, \text{NoC}, \text{NameC}, \text{WageC}} \\ (\sigma_{\text{Salary} > \text{WageC}}(\rho_{\text{NoC}, \text{NameC}, \text{WageC}, \text{AgeC} \leftarrow \text{No}, \text{Name}, \text{Wage}, \text{Age}} \text{Employee})) \\ \bowtie_{\text{NoC} = \text{Chief}} \\ (\text{Supervisor} \bowtie_{\text{Employee} = \text{Number}} ((\text{Employee}))) \end{array}$$

{ No: m, Name: n, Wage: s, NoC: c, NameC: nc, WageC: sc |
Employee(Number: m, Name: n, Age: a, Salary: s) \wedge
Supervisor(Chief:c, Employee:m) \wedge
Employee(Number: c, Name: nc, Age: ac, Salary: sc) \wedge s > sc}



Example (5)

- Return the chiefs' number and name having **all** employees earning more than 40

$$\begin{aligned} & \pi_{\text{Number}, \text{Name}} (\text{Employee} \bowtie_{\text{Number}=\text{Chief}} \\ & \quad (\pi_{\text{Chief}} (\text{Supervisor}) - \\ & \quad \pi_{\text{Chief}} (\text{Supervisor} \bowtie_{\text{Employee}=\text{Number}} \\ & \quad \sigma_{\text{Salary} \leq 40} (\text{Employee}))) \end{aligned}$$

{Number: c, Name: n |
Employee(Number: c, Name: n, Age: e, Salary: s) \wedge
Supervisor(Chief:c, Employee:m) \wedge
 $\neg \exists m'(\exists n'(\exists e'(\exists s'(\text{Employee}(No: m', \text{Name}: n', \text{Age}: e', \text{Wage}: s') \wedge
Supervisor(Chief:c, Employee:m') \wedge s' \leq 40))))}$



Recap

- De Morgan rules:

- $\neg(A \wedge B) = (\neg A) \vee (\neg B)$
- $\neg(A \vee B) = (\neg A) \wedge (\neg B)$

- Moreover:

- $\neg \exists x A = \forall x \neg A$
- $\neg \forall x A = \exists x \neg A$
- $\exists x \neg A = \neg \forall x A$

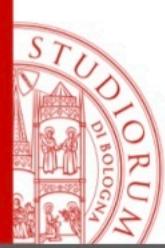


Which quantifier shall we use?

- Existential or universal? They are interchangeable by De Morgan:

{Number: c, Name: n |
Employee(Number: c, Name: n, Age: e, Salary: s) \wedge
Supervisor(Chief:c, Employee:m) \wedge
 $\neg \exists m'(\exists n'(\exists e'(\exists s'(\text{Employee}(No: m', Name: n', Age: e', Wage: s') \wedge
Supervisor(Chief:c, Employee:m') \wedge s' \leq 40))))}$

{Number: c, Name: n |
Employee(Number: c, Name: n, Age: e, Salary: s) \wedge
Supervisor(Chief:c, Employee:m) \wedge
 $\forall m'(\forall n'(\forall e'(\forall s'(\neg(\text{Employee}(No:m', Name:n', Age:e', Wage:s') \wedge
Supervisor(Chief:c, Employee:m')) \vee s' > 40))))}$



On Domain Relational Calculi

- Pros:
 - declarative
- Cons:
 - "verbose": so many variables!
 - Meaningless expressions:

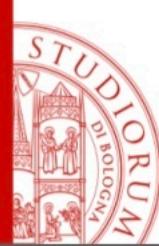
$\{ A: x \mid \neg R(A: x) \}$

$\{ A: x, B: y \mid R(A: x) \}$

$\{ A: x, B: y \mid R(A: x) \wedge y=y \}$

such expressions are "domain dependant" and we shall avoid them;

We cannot state such statements in relational algebra because it is domain independent.



Domain Calculus vs. Algebra

- Domain Relational Calculus (DRC) and Relational Algebra (RA) are "equivalent"
 - For each domain independent DRC expression, there exists an equivalent RA expression.
 - For each RA expression there exists an equivalent domain independent DRC expression



Tuples calculus with range declarations

- We want to overcome the domain calculus limitations:
 - We must reduce the number of variables; a good way to do so: we restrict the variables to the tuples, one variable for each tuple.
 - All the data values must come from the database
- The proposed calculus satisfies both needs



Tuple calculus with range declarations syntax

- The expressions have the following syntax:

$$\{ \text{TargetList} \mid \text{RangeList} \mid \text{Formula} \}$$

- *RangeList* shows the free variables in *Formula* specifying from which relation they come from.
- *TargetList* has elements like $Y: x.Z$ (or $x.Z$ or even $x.^*$)
- *Formula* has:
 - Comparison atoms $x.A \mathcal{R} c$, $x.A \mathcal{R} y.B$
 - Logical connectives
 - quantifiers associating a range over the variables

$$\exists x(R)(\dots) \quad \forall x(R)(\dots)$$



Example (0a)

- Return the employees' number, name, age and salary earning more than 40

$\sigma_{\text{Salary}>40}(\text{Employee})$

{ Number: m, Name: n, Age: e, Salary: s |
Employee(Number: m, Name: n, Age: e, Salary: s) \wedge s >
40 }

{ i.* | i(Employee) | i.Salary > 40 }



Example (0b)

- Return the employees' number, name and age

$\pi_{\text{Number, Name, Age}}(\text{Employee})$

{ Number: m, Name: n, Age: e |
Employee(Number: m, Name: n, Age: e, Salary: s)}

{ i.(Number,Name,Age) | i(Employee) | }



Example (1)

- Return the employees' number, name and age earning more than 40

$$\pi_{\text{Number, Name, Age}}(\sigma_{\text{Salary} > 40}(\text{Employee}))$$

{ Number: m, Name: n, Age: e |
Employee(Number: m, Name: n, Age: e, Salary: s) \wedge s >
40 }

{ i.(Number,Name,Age) | i(Employee) | i.Salary > 40 }



Example (2)

- Return the number identifying the employees' chiefs earning more than 40

{ Chief: c | Supervisor(Chief:c,Employee:m) \wedge
Employee(Number: m, Name: n, Age: e, Salary:
s) \wedge s > 40 }

{ s.Chief | i(Employee) , s(Supervisor) |
i.Number=s.Employee \wedge i.Salary > 40 }



Example (3)

- Return the chiefs' name and the salary having employees earning more than 40

```
{ NameC: nc, WageC: sc |  
Employee(Number: m, Name: n, Age: e, Salary: s) ∧ s >  
40 ∧  
Supervisor(Chief:c,Employee:m) ∧  
Employee(Number:c, Name:nc, Age:ec, Salary:sc) }
```

```
{ NameC,WageC: i'.(Name,Wage) |  
i'(Employee), s(Supervisor), i(Employee) |  
i'.Number=s.Chief ∧ i.Number=s.Employee ∧ i.Salary > 40 }
```



Example (4)

- Return the employees earning more money than their chiefs; for both such employees and chiefs return the number, name and salary

```
{ No: m, Name: n, Wage: s, NameC: nc, WageC: sc |  
  Employee(Number: m, Name: n, Age: e, Salary: s) ∧  
    Supervisor(Chef:c,Employee:m) ∧  
  Employee(Number: c, Name: nc, Age: ec, Salary: sc) ∧  
    s > sc}
```

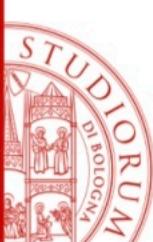
```
{ i.(Name,No,Wage), NameC,NoC,WageC: i'.(Name,No,Wage) |  
  i'(Employee), s(Supervisor), i(Employee) |  
  i'.Number=s.Chief ∧ i.Number=s.Employee ∧ i.Salary > i'.Salary }
```



Example (5)

- Return the chiefs' number and name having all employees earning more than 40

$$\{ \text{Number: } c, \text{ Name: } n \mid \\ \text{Employee}(\text{Number: } c, \text{ Name: } n, \text{ Age: } e, \text{ Salary: } s) \wedge \\ \text{Supervisor}(\text{Chief: } c, \text{ Employee: } m) \wedge \\ \neg \exists m' (\exists n' (\exists e' (\exists s' (\text{Employee}(\text{No: } m', \text{ Name: } n', \text{ Age: } e', \text{ Wage: } s') \wedge \\ \text{Supervisor}(\text{Chief: } c, \text{ Employee: } m') \wedge s' \leq 40))\}$$
$$\{ i.(\text{Number}, \text{ Name}) \mid s(\text{Supervisor}), i(\text{Employee}) \mid \\ i.\text{Number}=s.\text{Chief} \wedge \neg (\exists i'(\text{Employee})(\exists s'(\text{Supervisor}) \\ (s.\text{Chief}=s'.\text{Chief} \wedge s'.\text{Employee}=i'.\text{Number} \wedge i'.\text{Salary} \leq 40)))\}$$



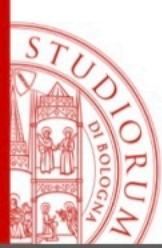
Remark

- Such calculus cannot express some important queries, such as unions:
$$R_1(AB) \cup R_2(AB)$$
- How could I express it through ranges? Is it possible with either one variable or two?
- On the other hand, we can express intersection and difference.
- For such reasons SQL (using this calculus) has an explicit union operator, while the intersection and difference operators do not appear in all SQL instances.



Tuple Calculus and Relational Algebra: Cons

- Both languages are basically equivalent: we can express a meaningful set of tuple operations with them.
- Some queries, potentially useful, cannot be expressed:
 - In both TC and RA we can only extract values, we cannot compute new values from them
 - Some interesting computations:
 - over each tuple (conversions, sums, differences, etc.)
 - over a set of tuples (summation, average, etc.)
- Such extensions are adopted in SQL, we will see them.
- Recursive queries, such as the **transitive closure**



Transitive Closure

Supervisor(Employee, Chief)

- For each employee, return all its superiors (e.g., its chief, the chief of its chief, and so on)

| Employee | Chief |
|----------|--------|
| Rossi | LUPI |
| Neri | Bruni |
| LUPI | Falchi |

| Employee | Superior |
|----------|----------|
| Rossi | LUPI |
| Neri | Bruni |
| LUPI | Falchi |
| Rossi | Falchi |

Transitive closure: how?

- We could use both joins with renaming in order to express such relations.
- But:

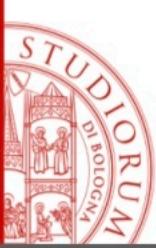
| Employee | Chief |
|----------|--------|
| Rossi | LUPI |
| Neri | Bruni |
| LUPI | Falchi |
| Falchi | Leoni |

| Employee | Superiors |
|----------|-----------|
| Rossi | LUPI |
| Neri | Bruni |
| LUPI | Falchi |
| Falchi | Leoni |
| Rossi | Falchi |
| LUPI | Leoni |
| Rossi | Leoni |



Transitive Closures are impossible

- In standard relational algebra we cannot express the transitive closure for each binary relation.
- In such languages, in order to express the transitive closure, we have each time to re-create a different expression:
 - How many join we would need?
 - There is no limit on the number of joins that are required!



Datalog

- A database oriented logical programming language, which ancestor is Prolog
- It uses two different types of predicates:
 - **extensional**: database's relations
 - **intensional**: correspond to “views”
- Intensional predicates are defined through **rules**



Datalog, Syntax

- Rules:

$$\textit{head} \leftarrow \textit{body}$$

- *head* is an atomic predicate (intensional)
- *body* is a list (conjunction) of atomic predicates
- Queries are specified by atomic predicates with a ? prefix



Preliminary Example

- Return the employees' number, name, age and salary being 30 years old

$$\{ \text{Number: } m, \text{Name: } n, \text{Age: } e, \text{Salary: } s \mid \\ \text{Employee}(\text{Number: } m, \text{Name: } n, \text{Age: } e, \text{Salary: } s) \wedge e = \\ 30 \}$$
$$? \text{Employee}(\text{Number: } m, \text{Name: } n, \text{Age: } 30, \text{Salary: } s)$$



Example (0a)

- Return the employees' number, name, age and salary earning more than 40

{ Number: m, Name: n, Age: e, Salary: s |
Employee(Number: m, Name: n, Age: e, Salary: s) \wedge s > 40 }

- We need an intensional predicate

RichE(Number: m, Name: n, Age: e, Salary: s) \leftarrow Employee(Number: m, Name: n, Age: e, Salary: s), s > 40

? RichE(Number: m, Name: n, Age: e, Salary: s)



Example (0b)

- Return the number, name and age for all the employees

$\pi_{\text{Number, Name, Age}}(\text{Employees})$

{ Number: m, Name: n, Age: e |
Employee(Number: m, Name: n, Age: e, Salary: s)}

PublicInfos(Number: m, Name: n, Age: e)
← Employee(Number: m, Name: n, Age: e, Salary: s)

? PublicInfos(Number: m, Name: n, Age: e)



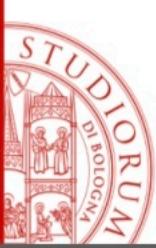
Example (2)

- Return the chief's number having employees earning more than 40

```
{ Chief: c | Supervisor(Chief:c,Employee:m) ∧  
Employee(Number: m, Name: n, Age: e, Salary: s) ∧ s > 40 }
```

```
ChiefsOfRichE (Chief:c) ←  
Employee(Number: m, Name: n, Age: e, Salary: s), s > 40,  
Supervisor (Chief:c,Employee:m)
```

```
? ChiefsOfRichE (Chief:c)
```



Example (5)

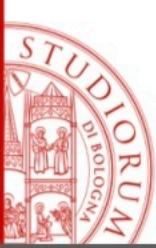
- Return the chiefs' number and name having employees earning ALL more than 40

- We need the negation

ChiefsOfNoRichE (Chief:c) ←
Supervisor (Chief:c,Employee:m),
Employee (Number: m, Name: n, Age: e, Salary: s),
 $s \leq 40$

ChiefsOfOnlyRichE(Number: c, Name: n) ←
Employee (Number: c, Name: n, Age: e, Salary: s) ,
Supervisor (Chief:c,Employee:m),
not ChiefsOfNoRichE (Chief:c)

? ChiefsOfOnlyRichE (Number: c, Name: n)

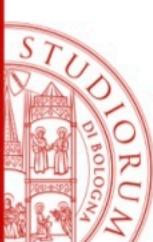


Example (6)

- For each employee, get all his/her chiefs.
- We need recursion

HighGrade (Employee: i, SuperChief: c) ←
Supervisor (Employee: i, Chief: c)

HighGrade (Employee: i, SuperChief: c) ←
Supervisor (Employee: i, Chief: c'),
HighGrade (Employee: c', SuperChief: c)



Datalog, semantics

- The definition of the recursive queries is quite tricky (in particular, the “negation” case)
- Expressive Power:
 - Non-recursive Datalog without negation is as expressive to the Calculus without negation and without universal quantifier
 - Non-recursive Datalog with negation is as expressive as calculus and algebra
 - We cannot compare Recursive Datalog without negation and Calculus
 - Recursive Datalog with negation is more expressive than calculus and algebra