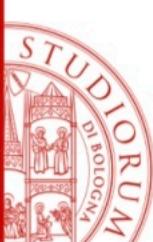




---

# **Databases**

## **Advanced SQL**



# Other Constraints: check

---

- “check” specifies constraints over tuples (and possibly, even more complex ones that aren’t always supported in all SQL implementations)

**CHECK ( *Predicate* )**



# Check, Example (1)

---

```
create table Employee
(
    Number integer primary key,
    Surname character(20),
    Name character(20),
    Sex character not null CHECK (Gender in ('M','F')) ,
    Salary integer CHECK (Salary >= 0) ,
    Supervisor integer,
    CHECK (Salary <= (select Salary
                        from Employee J
                       where Supervisor = J.Number) )
)
```

A **select** clause within the **CHECK** constraint is not fully supported in every SQL implementation.

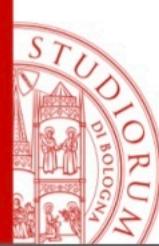


# Check, Example (2)

---

```
create table Employee
(
    Number character(6),
    Surname character(20),
    Name character(20),
    Sex character not null CHECK (Sex in ('M', 'F'))
    Salary integer,
    Withholding integer,
    Net integer,
    Supervisor character(6),
    check (Net = Salary - Withholding )
)
```

ok



# Check, Example (3)

---

<pre>insert into Employee values (1 , 'Doe', 'John', '', 100, 20, 80);</pre>	Sex is neither 'M' nor 'F'
<pre>insert into Employee values (2 , 'Smith', 'Jimmy', 'M', 100, 10, 80);</pre>	Net (80) is not 100-10
<pre>insert into Employee values (3 , 'Rossini', 'Luca', 'M', 70, 20, 50);</pre>	This update is allowed



# Other constraints: Assertion

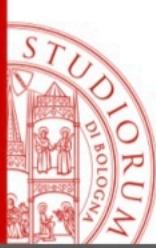
---

- Defines constraints at the schema level

```
create ASSERTION NameAs check ( Predicate )
```

```
create ASSERTION AtLeastOneEmployee  
    check (1 <= ( select count(*)  
                  from Employee ))
```

A **select** clause within the **check** constraint is not fully supported in every SQL implementation.



# View

---

```
create view ViewName [ ( AttList ) ] as
  SelectStatement
[ with [ local | cascaded ] check option ]
```

```
create view AdminEmployees
  (Name, Surname, Salary) as
select Name, Surname, Salary
from Employee
where Dept = 'Administration' and
  Salary > 10
```



# Querying a view

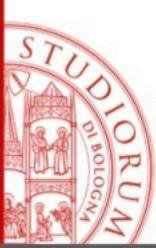
---

- Views could be queried like any other relation within the database. e.g.:

```
select * from AdminEmployees
```

is like performing the following query (and it is run as):

```
select Name, Surname, Salary  
from Employee  
where Dept = 'Administration' and  
Salary > 10
```



# View Update

---

- Usually, such updates are allowed for views defined on a single relation
- We can force the DB to perform some checks



# View Update: check option

---

```
create view PoorAdminEmployees as  
select *  
from AdminEmployees  
where Salary < 50  
with check option
```

- **check option** allows to update the view, but only if the inserted tuple belongs to the view (the user cannot have a salary greater than 50)



# Not allowed operation, Example

---

```
create view PoorAdminEmployees as
  select *
  from AdminEmployees
  where Salary < 50
with check option
```

```
update PoorAdminEmployees
  set Salary = 60
  where Name = 'Ann'
```



# Altering a View: local and cascaded

---

- **local** (in the case of views over views): the tuple update has to be performed only at the last level of the view
- **cascaded** (in the case of views over views): the tuple update has to be performed over all the underlying views and relations



# A wrong SQL query (1)

---

- Provide the average number of offices per department.
- Wrong query:  
`select avg(count(distinct Office))  
from Employee  
group by Dept`
- The query is wrong because the SQL syntax does not allow to nest aggregating operators



# A correct SQL query (1)

---

- Provide the average number of offices per department.

- Using an intermediate view:

```
create view DeptOffices(NameDept,OffNo) as  
select Dept, count(distinct Office)  
from Employee  
group by Dept;
```

```
select avg(OffNo)  
from DeptOffices
```



# A wrong SQL query (2)

---

- Provide the Dept. having the greatest employee's salary sum
- A solution wrong for some systems
  - `select Dept  
from Employee  
group by Dept  
having sum(Salary) >= all (select sum(Salary)  
from Employee  
group by Dept)`
- A nested **having** is not allowed in some SQL implementations

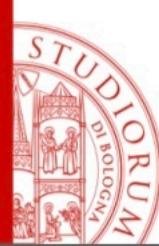


# A correct SQL query (2)

---

```
create view BudgetSalary(Dep,TotalSalary) as
select Dept, sum(Salary)
from Employee
group by Dept
```

```
select Dep
from BudgetSalary
where TotalSalary =(select max(TotalSalary)
                      from BudgetSalary)
```



# Recursive Queries (1)

- For each person provide his/her ancestors, having:  
**Fatherhood (Father, Child)**

Fatherhood	Father	Child
	Carl	Frank
	Louis	Olga
	Louis	Bob
	Frank	Alex
	Frank	Alfred



# Recursive Queries (2)

---

- We have to use recursion; in Datalog:

Ancestors (Ancestor: p, Descendant: f) ←  
Fatherhood (Father: p, Child: f)

Ancestors (Ancestor: a, Descendant: d) ←  
Fatherhood (Father: a, Child: f) ,  
Ancestors (Ancestor: f, Descendant: d)



# Recursive Queries in SQL:1999

---

```
with recursive Ancestors(Ancestor,Descendant) AS
( select Father, Son
  from Fatherhood
 union all
  select Ancestor, Son
  from Ancestors, Fatherhood
 where Descendant = Father
)
select *
```

**from Ancestors**

- **with** defines the **Ancestors** view, which is built up recursively using **Fatherhood**.



# Example

---

Return all John Doe's supervisors

**with recursive InCharge (No, Supervisor) AS**

**( select No, Supervisor**

**from Employee**

**union**

**select Employee.No, InCharge.Supervisor**

**from Employee, InCharge**

**where Employee.Supervisor = InCharge.No**

**)**

**select Name, Surname, InCharge.Supervisor**

**from Employee join InCharge**

**on (Employee.No = InCharge.No)**

**where Name = 'John' and Surname = 'Doe'**



# Scalar Functions (1)

---

- Functions at tuple level providing one single value per tuple
- Temporal
  - **current\_date** Returns the present date
  - **extract(*yearExpression*)**, extracts part of a date from a given expression (e.g. month, day, hour, etc.)
- Return the order year from the orders that were issued today

```
SELECT EXTRACT(YEAR from OrderDate) AS  
OrderYear,  
FROM Orders  
WHERE DATE(OrderDate)=current_date()
```



# Scalar Functions (2)

---

- String editing
  - **char\_length**, returns the string's length
  - **lower** , converts the string to lower case
- Casting
  - **Cast** allows the casting of a value into another domain;
- Conditional
  - ...



# Conditional Expression: coalesce (1)

**coalesce** takes several expressions as an input and returns the first of them which is not **NULL**.

- Given the following relation

Employee (Number, Dept, Mobile, PhoneHome)

For each employee return either a valid Mobile phone number, or its phone number.

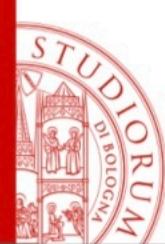
**Select** number, **coalesce**(Mobile, PhoneHome)  
**from** Employee



# Conditional Expression: coalesce (2)

**coalesce** could be used to provide a default value for replacing null values. In the following example, a NULL department is replaced with “None”

```
select Name, Surname, coalesce(Dept, 'None')  
from Employee
```



# Conditional Expression: **nullif**

---

**nullif** compares the first argument (e.g. an Attribute of a given relation) with the second one (e.g. a constant value). It returns NULL if the two arguments match, otherwise it returns the value from the first argument.

Example: Extract the surnames and departments to which the employees belong, returning a NULL value for the department when the Dept attribute has the value ‘Unknown’

```
select Surname, nullif(Dept, 'Unknown')  
from Employee
```



# Conditional Expression: case (1)

---

The **case** function allows to specify conditional structures, whose results depends on the evaluation of the tables' content.

It is used to provide an “if-then-else” kind of logic to the SQL language.



# Conditional Expression: case (2)

---

Evaluate the vehicle taxes by its type and registration year (since 1975).

```
select PlateNo,  
       (case Type  
             when 'Car' then 2.58 * Kwatt  
             when 'Moto' then (22.00 + 1.00 * Kwatt)  
             else null  
         end) as Tax  
from Vehicle  
where RegistrationYear > 1975
```



# Database Security

---

- SQL allows to grant for each user specific **privileges** (reading, writing, ...) for either the whole DB or part of it.
- **Privileges** could be **granted** to either relations, attributes, views or domains.
- There is at least one default admin (e.g. **\_system**) for which all the privileges are granted.
- Any user creating a specific “resource” automatically grants every privilege on it



# Privileges

---

A privilege is described by:

- a specific *resource*
- the user *granting* the privilege
- the user *to which* the privilege is granted
- a specific *operation*
- whether the user can *propagate* the privilege or not



# SQL's privileges

---

- **insert**: allows to insert new tuples
- **update**: allows to edit pre-existing tuples
- **delete**: allows to delete tuples
- **select**: allows to read a resource
- **references**: allows to define referential integrity constraints
- **usage**: allows to use a definition (e.g. a custom data type)



# Granting privileges

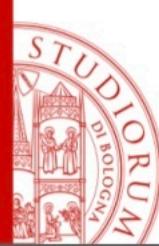
---

- granting privileges:

**grant < Privileges | all privileges > on  
Resource  
to Users [ with grant option ]**

- **grant option** allows the user to propagate his/her privilege to other users

**grant select on Department to Jack**



# Revoking privileges

---

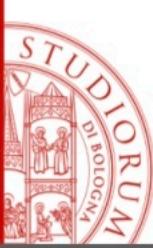
## ■ Revoking privileges

***revoke Privileges on Resource from Users***  
[ **restrict** | **cascade** ]

The privileges have to be revoked by the same user that granted them in the first place.

- **restrict** (default) the revoke does not involve other users that received the grant from the current user
- **cascade** the revoke is extended to the other users

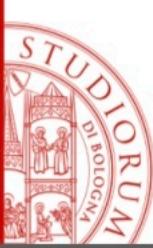
Be careful: **cascade** causes a “chain reaction”



# Privileges, discussion (1)

---

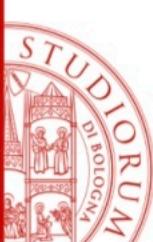
- The SQL implementation has to hide (without giving a clue) the part of the database that are not accessible to the user. E.g., either:
  - Table `Employees` doesn't exist, or
  - Table `Employees` does exist, but the user cannot access it,the user has to receive the same message from the system



# Privileges, discussion (2)

---

- We could use a view for showing only specific tuples to a user:
  - The view is defined with a selection predicate
  - The privilege is granted for the specific view



# Authorizations: RBAC (1)

---

SQL-3 has introduced RBAC for the first time (Role-Based Access Control, **RBAC**).

- Each role acts as a container of several privileges that can be granted via the `grant` command.
  
- At any moment, each user has both “individual” granted privileges, associated directly to him, and the privileges granted to its role through RBAC.



# Authorizations: RBAC(2)

---

- The following command creates a new role

**create role *Name***

- This other one grants *Name* to the current user

**set role *Name***



# RBAC, Example

---

- Grant **create table** to a specific user through the role Employee:

Create the new role:

```
create role Employee;
```

Grant the privilege to the previously defined role:

```
grant create table to Employee;
```

Grant the privilege to a specific user:

```
grant Employee to user1;
```

Revoke the previously granted privilege:

```
revoke create table from Employee;
```



# Transactions

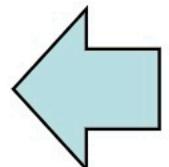
---

- A transaction is an executing program that forms a logical unit of database processing (atomic operation)
- Properties:
  - Atomicity
  - Consistency preservation
  - Isolation
  - Durability (permanency)

# Transactions are ... Atomic

---

- A transaction is an atomic unit of processing: it should either be performed or not performed at all.
- Money transfer from a bank account A to B: either the money is withdrawn from A and transferred to B, or no operation at all is performed.



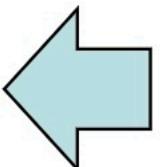
# Transactions are ... consistent

- A transaction should be consistency preserving: if it is completely executed from beginning to end with no interference, the database moves between two consistent states.
- During the execution of a transaction some violations may occur, but they cannot stay when the transaction ends: if in the end of a transaction there are still violations, the transaction must be cancelled (abort) completely.

# Transactions are ... isolated

---

- Even if many transactions are executed concurrently, the execution of the current transaction should not be interfered with others that are currently running.
  - E.g.: the case of two money transfers over the same bank account happening at the “same” time.

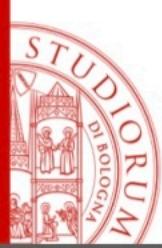




# Transactions have durable effects

---

- The changes applied to the database by a correct (**committed**) transaction must persist and must not be lost, even when serious faults (hardware or software) occur and in parallel execution.



# Transaction Support in SQL

---

- Once the connection to the database is established, the first transaction starts with **start transaction**, either before the first operation, or after the last one (the command is optional in most RDBMSs)
- After the first transaction it could terminate
  - **commit [work]**: the operations are saved within the database
  - **rollback [work]**: the operations are discarded and the DB returns to a previous safe state.
- Most RDBMSs have **autocommit**, where each statement is a different transaction.



# Transaction Support in SQL, Example

- Withdraw 10€ from the account 42177 and transfer it to the account 12202.

```
start transaction          (optional)
update BankAccount
    set Balance = Balance - 10
    where AccountNumber = 42177;
update BankAccount
    set Balance = Balance + 10
    where AccountNumber = 12202;
commit work;
```